

You have 2 free stories left this month. [Sign up](#) and get an extra one for free.

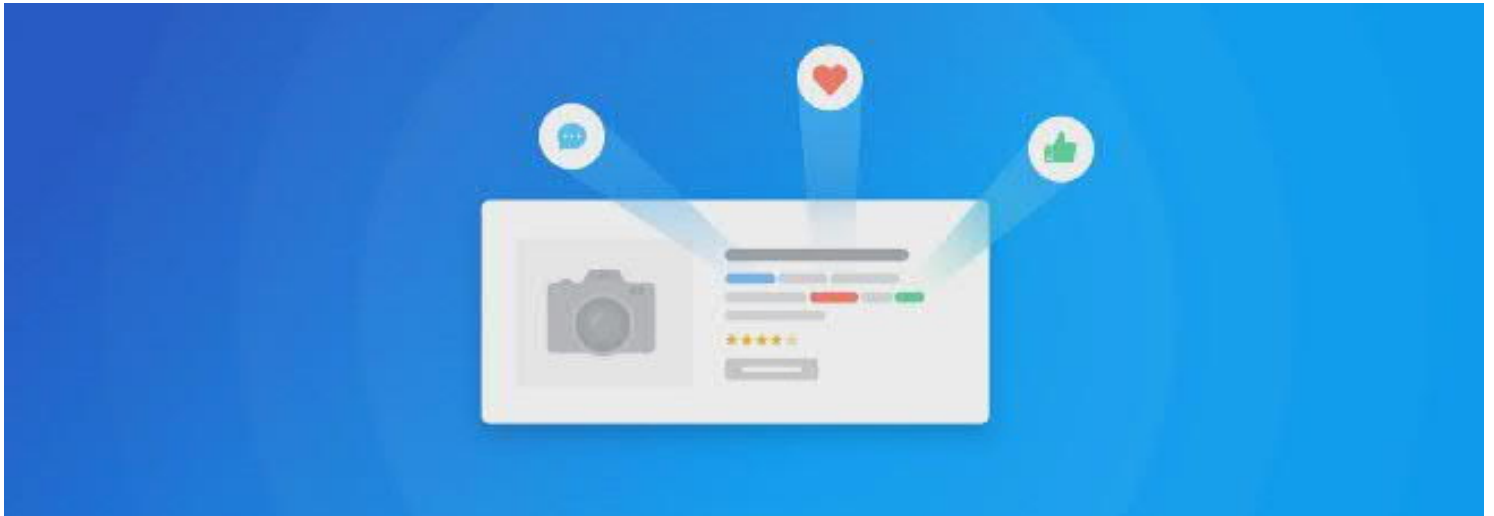
Sentiment Analysis using LSTM (Step-by-Step Tutorial)

Using PyTorch framework for Deep Learning



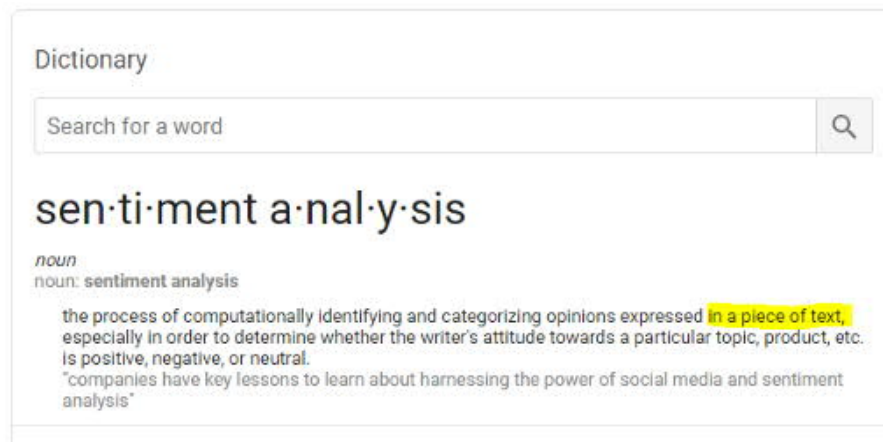
Samarth Agrawal [Follow](#)

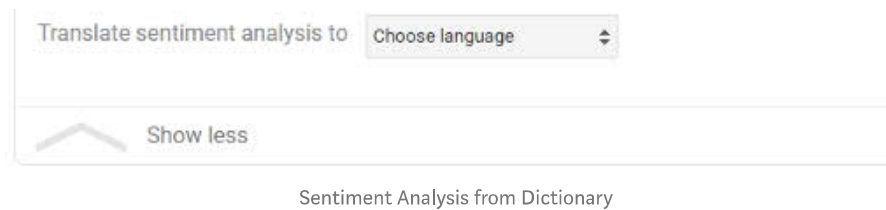
Feb 18, 2019 · 8 min read ★



Sentiment Analysis, Image by: [Monkeylearn](#)

What is Sentiment Analysis:





I think this result from google dictionary gives a very succinct definition. I don't have to re-emphasize how important sentiment analysis has become. So, here we will build a classifier on IMDB movie dataset using a Deep Learning technique called RNN.

I'm outlining a step-by-step process for how Recurrent Neural Networks (RNN) can be implemented using Long Short Term Memory (LSTM) architecture:

1. Load in and visualize the data
2. Data Processing — convert to lower case
3. Data Processing — Remove punctuation
4. Data Processing — Create list of reviews
5. Tokenize — Create Vocab to Int mapping dictionary
6. Tokenize — Encode the words
7. Tokenize — Encode the labels
8. Analyze Reviews Length
9. Removing Outliers — Getting rid of extremely long or short reviews
10. Padding / Truncating the remaining data
11. Training, Validation, Test Dataset Split
12. Dataloaders and Batching
13. Define the LSTM Network Architecture
14. Define the Model Class
15. Training the Network
16. Testing (on Test data and User- generated data)

. . .

1) Load in and visualize the data

We are using IMDB movies review dataset. If it is stored in your machine in a txt file then we just load it in

```
# read data from text files
with open('data/reviews.txt', 'r') as f:
    reviews = f.read()
with open('data/labels.txt', 'r') as f:
    labels = f.read()

print(reviews[:50])
print()
print(labels[:26])

--- Output ---

bromwell high is a cartoon comedy . it ran at the same time as some
other programs about school life  such as  teachers  . my  years

positive
negative
positive
```

2) Data Processing — convert to lower case

```
reviews = reviews.lower()
```

3) Data Processing — remove punctuation

```
from string import punctuation
print(punctuation)

--- Output ---

!"#$%&'()*+,-./:;<=>?@[\\]^_`{|}~
```

We saw all the punctuation symbols predefined in python. To get rid of all these punctuation we will simply use

```
all_text = ''.join([c for c in reviews if c not in punctuation])
```

4) Data Processing — create list of reviews

We have got all the strings in one huge string. Now we will separate out individual reviews and store them as individual list elements. Like,
[review_1, review_2, review_3..... review_n]

```
reviews_split = all_text.split('\n')
print ('Number of reviews :', len(reviews_split))
```

Number of reviews : 25001

5) Tokenize — Create Vocab to Int mapping dictionary

In most of the NLP tasks, you will create an index mapping dictionary in such a way that your frequently occurring words are assigned lower indexes. One of the most common way of doing this is to use `Counter` method from `Collections` library.

```
from collections import Counter

all_text2 = ' '.join(reviews_split)
# create a list of words
words = all_text2.split()

# Count all the words using Counter Method
count_words = Counter(words)

total_words = len(words)
sorted_words = count_words.most_common(total_words)
```

Let's have a look at these objects we have created

```
print (count_words)

--- Output ---

Counter({'the': 336713, 'and': 164107, 'a': 163009, 'of': 145864
```

In order to create a vocab to int mapping dictionary, you would simply do this

```
vocab_to_int = {w:i for i, (w,c) in enumerate(sorted_words)}
```

There is a small trick here, in this mapping index will start from 0 i.e. mapping of 'the' will be 0. But later on we are going to do padding for shorter reviews and conventional choice for padding is 0. So we need to start this indexing from 1

```
vocab_to_int = {w:i+1 for i, (w,c) in enumerate(sorted_words)}
```

Let's have a look at this mapping dictionary. We can see that mapping for 'the' is 1 now

```
print (vocab_to_int)

--- Output ---

{'the': 1, 'and': 2, 'a': 3, 'of': 4,
```

6) Tokenize — Encode the words

So far we have created a) list of reviews and b) index mapping dictionary using vocab from all our reviews. All this was to create an encoding of reviews (replace words in our reviews by integers)

```
reviews_int = []
for review in reviews_split:
    r = [vocab_to_int[w] for w in review.split()]
    reviews_int.append(r)
print (reviews_int[0:3])

--- Output ---

[[21025, 308, 6, 3, 1050, 207, 8, 2138, 32, 1, 171, 57, 15, 49, 81,
5785, 44, 382, 110, 140, 15, .....], [5194, 60, 154, 9, 1, 4975,
5852, 475, 71, 5, 260, 12, 21025, 308, 13, 1978, 6, 74, 2395, 5, 613,
73, 6, 5194, 1, 24103, 5, ....], [1983, 10166, 1, 5786, 1499, 36, 51,
66, 204, 145, 67, 1199, 5194.....]]
```

Note: what we have created now is a list of lists. Each individual review is a list of integer values and all of them are stored in one huge list

7) Tokenize — Encode the labels

This is simple because we only have 2 output labels. So, we will just label 'positive' as 1 and 'negative' as 0

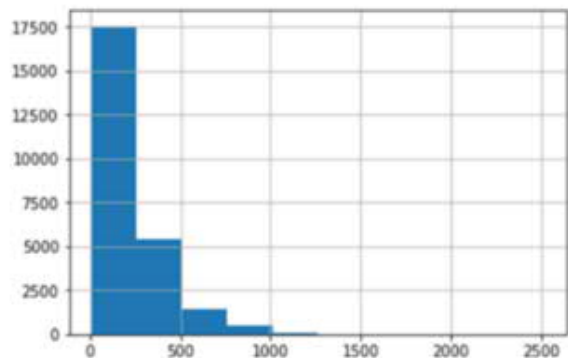
```
encoded_labels = [1 if label == 'positive' else 0 for label in
labels_split]
encoded_labels = np.array(encoded_labels)
```

8) Analyze Reviews Length

```
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline

reviews_len = [len(x) for x in reviews_int]
pd.Series(reviews_len).hist()
plt.show()

pd.Series(reviews_len).describe()
```



```
count    25001.000000
mean      240.798208
std       179.020628
min         0.000000
25%       130.000000
50%       179.000000
75%       293.000000
max      2514.000000
dtype: float64
```

Review Length Analysis

Observations : a) Mean review length = 240 b) Some reviews are of 0 length. Keeping this review won't make any sense for our analysis c) Most of the reviews less than 500 words or more d) There are quite a few reviews that are extremely long, we can manually investigate them to check whether we need to include or exclude them from our analysis

9) Removing Outliers — Getting rid of extremely long or short reviews

```
reviews_int = [ reviews_int[i] for i, l in enumerate(reviews_len) if l>0 ]
encoded_labels = [ encoded_labels[i] for i, l in enumerate(reviews_len) if l> 0 ]
```

10) Padding / Truncating the remaining data

To deal with both short and long reviews, we will pad or truncate all our reviews to a specific length. We define this length by **Sequence Length**. This sequence length is same as number of time steps for LSTM layer.

For reviews shorter than `seq_length`, we will pad with 0s. For reviews longer than `seq_length` we will truncate them to the first `seq_length` words.

```
def pad_features(reviews_int, seq_length):
    ''' Return features of review_ints, where each review is padded
    with 0's or truncated to the input seq_length.
    '''
    features = np.zeros((len(reviews_int), seq_length), dtype = int)

    for i, review in enumerate(reviews_int):
        review_len = len(review)

        if review_len <= seq_length:
            zeroes = list(np.zeros(seq_length-review_len))
            new = zeroes+review

        elif review_len > seq_length:
            new = review[0:seq_length]

        features[i,:] = np.array(new)

    return features
```

Note: We are creating/maintaining a 2D array structure as we created for `reviews_int` . Output will look like this

```
print (features[:10,:])
```

```
[[ 0  0  0 ...  8 215 23]
 [ 0  0  0 ... 29 108 3324]
 [22382 42 46418 ... 483 17 3]
 ...
 [ 0  0  0 ... 59 429 1776]
 [ 0  0  0 ... 55 201 18]
 [ 0  0  0 ... 18 17 1191]]
```

11) Training, Validation, Test Dataset Split

Once we have got our data in nice shape, we will split it into training, validation and test sets

train= 80% | valid = 10% | test = 10%

```
split_frac = 0.8
train_x = features[0:int(split_frac*len_feat)]
train_y = encoded_labels[0:int(split_frac*len_feat)]

remaining_x = features[int(split_frac*len_feat):]
remaining_y = encoded_labels[int(split_frac*len_feat):]

valid_x = remaining_x[0:int(len(remaining_x)*0.5)]
valid_y = remaining_y[0:int(len(remaining_y)*0.5)]

test_x = remaining_x[int(len(remaining_x)*0.5):]
test_y = remaining_y[int(len(remaining_y)*0.5):]
```

12) Dataloaders and Batching

After creating our training, test and validation data. Next step is to create dataloaders for this data. We can use generator function for batching our data into batches instead we will use a TensorDataset. This is one of a very useful utility in PyTorch for using our data with DataLoaders with exact same ease as of torchvision datasets

```
import torch
from torch.utils.data import DataLoader, TensorDataset

# create Tensor datasets
train_data = TensorDataset(torch.from_numpy(train_x),
                             torch.from_numpy(train_y))
valid_data = TensorDataset(torch.from_numpy(valid_x),
                             torch.from_numpy(valid_y))
test_data = TensorDataset(torch.from_numpy(test_x),
                             torch.from_numpy(test_y))

# dataloaders
batch_size = 50

# make sure to SHUFFLE your data
train_loader = DataLoader(train_data, shuffle=True,
                           batch_size=batch_size)
valid_loader = DataLoader(valid_data, shuffle=True,
                           batch_size=batch_size)
test_loader = DataLoader(test_data, shuffle=True,
                          batch_size=batch_size)
```

In order to obtain one batch of training data for visualization purpose we will create a data iterator

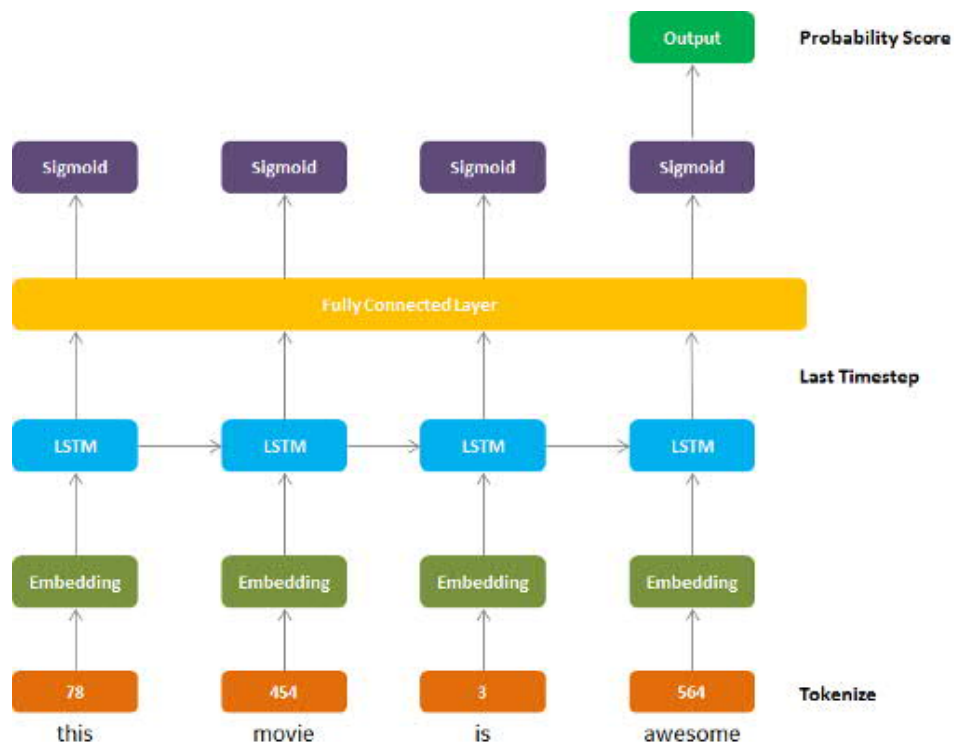
```
# obtain one batch of training data
dataiter = iter(train_loader)
sample_x, sample_y = dataiter.next()

print('Sample input size: ', sample_x.size()) # batch_size,
seq_length
print('Sample input: \n', sample_x)
print()
print('Sample label size: ', sample_y.size()) # batch_size
print('Sample label: \n', sample_y)
```

```
Sample input size: torch.Size([50, 200])  
Sample input:  
tensor([[ 24,    65, 1011, ..., 4692,     2, 1999],  
        [ 72,   87,   14, ..., 6667,    31,    4],  
        [  1,  240,    9, ..., 1940,     7,    7],  
        ...,  
        [ 11,   18,   14, ...,  151,    23,   786],  
        [  0,    0,    0, ...,   21,   397,   243],  
        [  0,    0,    0, ...,  129, 1314, 3622]], dtype=torch.int32)  
  
Sample label size: torch.Size([50])  
Sample label:  
tensor([0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1,  
        1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0,  
        0, 0], dtype=torch.int32)
```


Here, 50 is the batch size and 200 is the sequence length that we have defined. Now our data prep step is complete and next we will look at the LSTM network architecture for start building our model

13) Define the LSTM Network Architecture



LSTM Architecture for Sentiment Analysis

The layers are as follows:

0. Tokenize : This is not a layer for LSTM network but a mandatory step of converting our words into tokens (integers)
1. Embedding Layer: that converts our word tokens (integers) into embedding of specific size
2. LSTM Layer: defined by hidden state dims and number of layers
3. Fully Connected Layer: that maps output of LSTM layer to a desired output size
4. Sigmoid Activation Layer: that turns all output values in a value between 0 and 1
5. Output: Sigmoid output from the last timestep is considered as the final output of this network

14) Define the Model Class

```

1  import torch.nn as nn
2
3  class SentimentLSTM(nn.Module):
4      """
5      The RNN model that will be used to perform Sentiment analysis.
6      """
7
8      def __init__(self, vocab_size, output_size, embedding_dim, hidden_dim, n_layers, drop_prob=0
9          """
10         Initialize the model by setting up the layers.
11         """
12         super().__init__()
13
14         self.output_size = output_size
15         self.n_layers = n_layers
16         self.hidden_dim = hidden_dim
17
18         # embedding and LSTM layers
19         self.embedding = nn.Embedding(vocab_size, embedding_dim)
20         self.lstm = nn.LSTM(embedding_dim, hidden_dim, n_layers,
21                             dropout=drop_prob, batch_first=True)
22
23         # dropout layer
24         self.dropout = nn.Dropout(0.3)
25
26         # linear and sigmoid layers
27         self.fc = nn.Linear(hidden_dim, output_size)
28         self.sig = nn.Sigmoid()
29
30
31     def forward(self, x, hidden):
32         """
33         Perform a forward pass of our model on some input and hidden state.
34         """
35         batch_size = x.size(0)
36
37         # embeddings and lstm_out
38         embeds = self.embedding(x)
39         lstm_out, hidden = self.lstm(embeds, hidden)
40
41         # stack up lstm outputs
42         lstm_out = lstm_out.contiguous().view(-1, self.hidden_dim)
43
44         # dropout and fully-connected layer
45         out = self.dropout(lstm_out)
46         out = self.fc(out)
47         # sigmoid function
48         sig_out = self.sig(out)
49
50         # reshape to be batch_size first
51         sig_out = sig_out.view(batch_size, -1)
52         sig_out = sig_out[:, -1] # get last batch of labels
53
54         # return last sigmoid output and hidden state
55         return sig_out, hidden
56
57
58     def init_hidden(self, batch_size):
59         ''' Initializes hidden state '''
60         # Create two new tensors with sizes n_layers x batch_size x hidden_dim,

```

```

61         # initialized to zero, for hidden state and cell state of LSTM
62         weight = next(self.parameters()).data
63
64         if (train_on_gpu):
65             hidden = (weight.new(self.n_layers, batch_size, self.hidden_dim).zero_().cuda(),
66                       weight.new(self.n_layers, batch_size, self.hidden_dim).zero_().cuda())
67         else:
68             hidden = (weight.new(self.n_layers, batch_size, self.hidden_dim).zero_(),
69                       ... ..)

```

```

# Instantiate the model w/ hyperparams
vocab_size = len(vocab_to_int)+1 # +1 for the 0 padding
output_size = 1
embedding_dim = 400
hidden_dim = 256
n_layers = 2

net = SentimentLSTM(vocab_size, output_size, embedding_dim,
                    hidden_dim, n_layers)

print(net)

SentimentLSTM(
  (embedding): Embedding(74073, 400)
  (lstm): LSTM(400, 256, num_layers=2, batch_first=True, dropout=0.5)
  (dropout): Dropout(p=0.3)
  (fc): Linear(in_features=256, out_features=1, bias=True)
  (sig): Sigmoid()
)

```

- Training Loop

Most of the code in training loop is pretty standard Deep Learning training code that you might see often in all the implementations that's using PyTorch framework.

```

1  # loss and optimization functions
2  lr=0.001
3
4  criterion = nn.BCELoss()
5  optimizer = torch.optim.Adam(net.parameters(), lr=lr)
6
7
8  # training params
9
10 epochs = 4 # 3-4 is approx where I noticed the validation loss stop decreasing
11
12 counter = 0
13 print_every = 100
14 clip=5 # gradient clipping
15
16 # move model to GPU, if available
17 if(train_on_gpu):

```

```

18     net.cuda()
19
20 net.train()
21 # train for some number of epochs
22 for e in range(epochs):
23     # initialize hidden state
24     h = net.init_hidden(batch_size)
25
26     # batch loop
27     for inputs, labels in train_loader:
28         counter += 1
29
30         if(train_on_gpu):
31             inputs, labels = inputs.cuda(), labels.cuda()
32
33         # Creating new variables for the hidden state, otherwise
34         # we'd backprop through the entire training history
35         h = tuple([each.data for each in h])
36
37         # zero accumulated gradients
38         net.zero_grad()
39
40         # get the output from the model
41         inputs = inputs.type(torch.LongTensor)
42         output, h = net(inputs, h)
43
44         # calculate the loss and perform backprop
45         loss = criterion(output.squeeze(), labels.float())
46         loss.backward()
47         # `clip_grad_norm` helps prevent the exploding gradient problem in RNNs / LSTMs.
48         nn.utils.clip_grad_norm_(net.parameters(), clip)
49         optimizer.step()
50
51     # loss stats
52     if counter % print_every == 0:
53         # Get validation loss
54         val_h = net.init_hidden(batch_size)
55         val_losses = []
56         net.eval()
57         for inputs, labels in valid_loader:
58
59             # Creating new variables for the hidden state, otherwise
60             # we'd backprop through the entire training history
61             val_h = tuple([each.data for each in val_h])
62
63             if(train_on_gpu):
64                 inputs, labels = inputs.cuda(), labels.cuda()
65
66             inputs = inputs.type(torch.LongTensor)
67             output, val_h = net(inputs, val_h)
68             val_loss = criterion(output.squeeze(), labels.float())
69
70             val_losses.append(val_loss.item())
71
72 net.train()

```

[

- On Test Data

```

1  # Get test data loss and accuracy
2
3  test_losses = [] # track loss
4  num_correct = 0
5
6  # init hidden state
7  h = net.init_hidden(batch_size)
8
9  net.eval()
10 # iterate over test data
11 for inputs, labels in test_loader:
12
13     # Creating new variables for the hidden state, otherwise
14     # we'd backprop through the entire training history
15     h = tuple([each.data for each in h])
16
17     if(train_on_gpu):
18         inputs, labels = inputs.cuda(), labels.cuda()
19
20     # get predicted outputs
21     inputs = inputs.type(torch.LongTensor)
22     output, h = net(inputs, h)
23
24     # calculate loss
25     test_loss = criterion(output.squeeze(), labels.float())
26     test_losses.append(test_loss.item())
27
28     # convert output probabilities to predicted class (0 or 1)
29     pred = torch.round(output.squeeze()) # rounds to the nearest integer
30
31     # compare predictions to true label
32     correct_tensor = pred.eq(labels.float().view_as(pred))
33     correct = np.squeeze(correct_tensor.numpy()) if not train_on_gpu else np.squeeze(correct_tensor.cpu().numpy())
34     num_correct += np.sum(correct)
35
36
37 # -- stats! -- ##
38 # avg test loss
39 print("Test loss: {:.3f}".format(np.mean(test_losses)))
40
41 # accuracy over all test data
42 test_acc = num_correct/len(test_loader.dataset)

```

First, we will define a `tokenize` function that will take care of pre-processing steps and then we will create a `predict` function that will give us the final output after parsing the user provided review.

```

1  from string import punctuation
2
3  def tokenize_review(test_review):
4      test_review = test_review.lower() # lowercase
5      # get rid of punctuation
6      test_text = ''.join([c for c in test_review if c not in punctuation])
7
8      # splitting by spaces
9      test_words = test_text.split()
10

```

```
11     # tokens
12     test_ints = []
13     test_ints.append([vocab_to_int[word] for word in test_words])
14
15     return test_ints
16
17 # test code and generate tokenized review
18 test_ints = tokenize_review(test_review_neg)
19 print(test_ints)
20
21
22 # test sequence padding
23 seq_length=200
24 features = pad_features(test_ints, seq_length)
25
26 print(features)
27
28
29 # test conversion to tensor and pass into your model
30 feature_tensor = torch.from_numpy(features)
31 print(feature_tensor.size())
32
33
34 def predict(net, test_review, sequence_length=200):
35
36     net.eval()
37
38     # tokenize review
39     test_ints = tokenize_review(test_review)
40
41     # pad tokenized sequence
42     seq_length=sequence_length
43     features = pad_features(test_ints, seq_length)
44
45     # convert to tensor to pass into your model
46     feature_tensor = torch.from_numpy(features)
47
48     batch_size = feature_tensor.size(0)
49
50     # initialize hidden state
51     h = net.init_hidden(batch_size)
52
53     if(train_on_gpu):
54         feature_tensor = feature_tensor.cuda()
55
56     # get the output from the model
57     output, h = net(feature_tensor, h)
58
59     # convert output probabilities to predicted class (0 or 1)
60     pred = torch.round(output.squeeze())
61     # printing output value, before rounding
62     print('Prediction value, pre-rounding: {:.6f}'.format(output.item()))
63
64     # print custom response
65     if(pred.item() != 1):
```

```
test_review = 'This movie had the best acting and the dialogue was so
good. I loved it.'

seq_length=200 # good to use the length that was trained on

predict(net, test_review_neg, seq_length)
```

Positive review detected

Closing Remarks:

- I have tried to detail out the process involved in building a Sentiment Analysis classifier based on LSTM architecture using PyTorch framework.
- Please feel free to write your thoughts / suggestions / feedbacks

Update: Another article to give you a microscopic view of what happens within the layers. [Read here](#)

Sign up for The Daily Pick

By Towards Data Science

Hands-on real-world examples, research, tutorials, and cutting-edge techniques delivered Monday to Thursday. Make learning your daily ritual. [Take a look](#)

 Get this newsletter

Create a free Medium account to get The Daily Pick in your inbox.

Machine Learning

NLP

Sentiment Analysis

Deep Learning

Lstm

Discover Medium

Welcome to a place where words matter. On Medium, smart voices and original ideas take center stage - with no ads in sight. [Watch](#)

Make Medium yours

Follow all the topics you care about, and we'll deliver the best stories for you to your homepage and inbox. [Explore](#)

Become a member

Get unlimited access to the best stories on Medium — and support writers while you're at it. Just \$5/month. [Upgrade](#)

Medium

[About](#)

[Help](#)

[Legal](#)